

# Verteilte Systeme

...für C++ Programmierer

Threads 2

by

Dr. Günter Kolousek

# Threads und 'callable types'

- ▶ Thread kann jeden beliebigen 'callable' Typ bei Initialisierung als Parameter bekommen
- ▶ Dies sind:
  - ▶ Normale Funktionen (per Funktionsname, per Pointer auf Funktion)
  - ▶ Lambda-Ausdrücke
  - ▶ `std::function` - Objekte (aus `<functional>`)
  - ▶ Klassen mit überladenen Operator `operator()`.

# Lambda-Thread

```
#include <iostream> // lambdathread.cpp
#include <thread>
using namespace std;

int main() {
    thread t{[]() {
        cout << "lambda thread" << endl; }};
    cout << "main thread" << endl;
    t.join();
}
```

# Thread mit Wertparameter

```
#include <iostream> // valparthread.cpp
#include <thread>
using namespace std;

int main() {
    auto outfunc = [](string const& msg) {
        cout << "lambda " << msg << endl;
    };
    thread t{outfunc,
             "thread"};
    cout << "main thread" << endl;
    t.join();
}
```

```
main threadlambda
thread
```

# Thread mit Wertparameter - 2

```
#include <iostream>    // valparthread2.cpp
#include <thread>
using namespace std;
void f() {
    char buffer[1024]{"very long string..."};
    auto outfunc{[](string msg) {
        cout << "lambda " << msg << endl; }};
    thread t{outfunc, buffer};
    t.detach();
}
int main() {
    f();    this_thread::sleep_for(10ms);
}
unsicher!!!
```

# Thread mit Wertparameter - 2

```
#include <iostream>    // valparthread2.cpp
#include <thread>
using namespace std;
void f() {
    char buffer[1024]{"very long string..."};
    auto outfunc{[](string msg) {
        cout << "lambda " << msg << endl; }};
    thread t{outfunc, buffer};
    t.detach();
}
int main() {
    f();    this_thread::sleep_for(10ms);
}
```

unsicher!!! buffer wird als char\* übergeben (per value) und im Kontext des Threads wird ein string Objekt erzeugt...

# Thread mit Wertparameter - 2

```
#include <iostream>    // valparthread2.cpp
#include <thread>
using namespace std;
void f() {
    char buffer[1024]{"very long string..."};
    auto outfunc{[](string msg) {
        cout << "lambda " << msg << endl; }};
    thread t{outfunc, buffer};
    t.detach();
}
int main() {
    f();    this_thread::sleep_for(10ms);
}
```

unsicher!!! buffer wird als char\* übergeben (per value) und im Kontext des Threads wird ein string Objekt erzeugt... besser:

```
thread t{outfunc, string{buffer}};
```

# Thread mit Referenzparameter

```
#include <iostream>    // refparthread.cpp
#include <thread>
#include <functional>    // ref
using namespace std;

int main() {
    int n{};
    auto incrfunc{[](int& n) { ++n; }};
    // use std::ref, otherwise per-value!
    // generates 'reference_wrapper'
    thread t{incrfunc, ref(n)};
    t.join();
    cout << n << endl; }
```

1

Aber Achtung, wenn Speicherobjekt nicht mehr vorhanden...



# Thread mit Referenzparameter – 2

```
#include <iostream> // refparthread2.cpp
#include <thread>
#include <functional> // ref
using namespace std;
using namespace literals;

struct Distance {
    double len{};
    ~Distance() { cout << "dstor" << endl; }
};
```

# Thread mit Referenzparameter – 3

```
int main() {
    Distance* pd;
    {
        Distance d{};
        pd = &d;
        auto incrfunc{[](Distance& d) {
            this_thread::sleep_for(1s); ++d.len; }}
        thread t(incrfunc, ref(d)); // be careful!
        t.detach();
    }
    this_thread::sleep_for(2s);
    cout << "trying to access pd->len..." << endl;
}
```

# Thread mit Referenzparameter – 3

```
int main() {  
    Distance* pd;  
    {  
        Distance d{};  
        pd = &d;  
        auto incrfunc{[](Distance& d) {  
            this_thread::sleep_for(1s); ++d.len; }}  
        thread t(incrfunc, ref(d)); // be careful!  
        t.detach();  
    }  
    this_thread::sleep_for(2s);  
    cout << "trying to access pd->len..." << endl;  
}
```

dstor  
trying to access pd->len...

# Thread - moveable Argument

```
#include <iostream>    // movparthread.cpp
#include <thread>
using namespace std;
struct Ressource {
    int value{};
    int id{};
    Ressource(int value_) : value{value_} {}
    Ressource(Ressource&& o) {
        value = o.value;
        id = o.id + 1;
        cout<< "from " << o.id<< " to " << id<< endl;
        o.value = 0;
    }
    ~Ressource() { cout << "dtor: " << id <<
                  " value: " << value << endl; }
};
```

# Thread - moveable Argument – 2

```
int main() {  
    thread t1{[] (Ressource&& res) { cout <<  
        "t1: " << res.value << endl; },  
        Ressource{42}};  
    t1.join();  
}
```

from 0 to 1

from 1 to 2

dtor of id 1 value: 0

dtor of id 0 value: 0

t1: 42

dtor of id 2 value: 42

# Thread ID

```
#include <iostream> // threadid.cpp
#include <thread>
using namespace std;
using namespace std::literals;
int main() {
    thread t{[]() { this_thread::sleep_for(1s); }};
    // type std::thread::id
    // comparable (<,<=,==,...) & storable in map
    cout << this_thread::get_id() << ' ';
    cout << t.get_id() << endl;
    t.join();
    thread t2; // no hw thread associated
    cout << t2.get_id() << endl; }
```

140524781053760 140524781049600

thread::id of a non-executing thread

# Thread-Affinity (Linux)

```
#include <iostream>    // affinity.cpp
#include <vector>
#include <thread>
#include <mutex>
#include <sched.h>    // sched_getcpu()
using namespace std;
using namespace std::literals;
```

# Thread-Affinity (Linux) – 2

```
int main(int argc, const char** argv) {
    constexpr unsigned num_threads = 4;
    mutex mtx;
    vector<thread> threads(num_threads);
    for (unsigned i = 0; i < num_threads; ++i) {
        threads[i] = thread([&mtx, i] {
            while (true) {
                {
                    lock_guard<mutex> iolock(mtx);
                    cout<<"Thread #"<<i<<": on CPU "
                        <<sched_getcpu()<<"\n";
                }
                this_thread::sleep_for(900ms);
            } });
    }
    for (auto& t : threads) t.join();
}
```



# Thread-Affinity (Linux) – 3

Mögliche Ausgabe des Programmes affinity

```
Thread #1: on CPU 2  
Thread #0: on CPU 0  
Thread #2: on CPU 1  
Thread #3: on CPU 0  
Thread #1: on CPU 2  
Thread #0: on CPU 3  
Thread #3: on CPU 1  
Thread #2: on CPU 2  
...
```

# Thread-Affinity (Linux) – 4

Nochmals mit nur den Kernen 2 und 3:

```
$ taskset -c 2,3 affinity
```

Thread #0: on CPU 2

Thread #1: on CPU 2

Thread #2: on CPU 3

Thread #3: on CPU 2

Thread #0: on CPU 2

Thread #2: on CPU 3

Thread #1: on CPU 3

Thread #3: on CPU 2

...

# Setzen der Thread-Affinity (Linux)

```
#include <pthread.h>

#include <iostream> // setaffinity.cpp
#include <vector>
#include <thread>
#include <mutex>
#include <sched.h> //

using namespace std;
using namespace std::literals;
```

# Setzen der Thread-Affinity – 2

```
int main(int argc, const char** argv) {
    constexpr unsigned num_threads = 4;
    mutex mtx;
    vector<thread> threads(num_threads);
    for (unsigned i{}; i < num_threads; ++i) {
        threads[i] = thread([&mtx, i] {
            this_thread::sleep_for(20ms);
            while (true) {
                {
                    lock_guard<mutex> iolock(mtx);
                    cout << "Thread #" << i
                        << ": on CPU "
                        << sched_getcpu() << "\n";
                }
                this_thread::sleep_for(900ms);
            }
        });
    }
}
```

# Setzen der Thread-Affinity – 3

```
// cpu_set_t represents a set of CPUs.  
// Clear it and mark only CPU i as set.  
cpu_set_t cpuset;  
CPU_ZERO(&cpuset);  
CPU_SET(i, &cpuset);  
int rc = pthread_setaffinity_np(  
    threads[i].native_handle(),  
    sizeof(cpu_set_t), &cpuset);  
if (rc != 0) {  
    cerr <<  
        "Error calling pthread_setaffinity_np:  
    << rc << "\n";  
}  
}  
for (auto& t : threads) t.join();  
}
```

# Setzen der Thread-Affinity – 4

Mögliche Ausgabe:

Thread #0: on CPU 0

Thread #2: on CPU 2

Thread #3: on CPU 3

Thread #1: on CPU 1

Thread #0: on CPU 0

Thread #2: on CPU 2

Thread #1: on CPU 1

Thread #3: on CPU 3

...